Task 1

```
CREATE PARTITION FUNCTION monthlyConsumptionYears(CHAR(4))
AS RANGE LEFT FOR VALUES (2020, 2021, 2022, 2023);

CREATE PARTITION SCHEME monthlyConsumptionYearsPartition
AS PARTITION monthlyConsumptionYears
TO ([PRIMARY], [PRIMARY], [PRIMARY], [PRIMARY], [PRIMARY]);

ALTER TABLE Energy.MonthlyConsumption DROP CONSTRAINT
PK_MonthlyConsumption;

CREATE CLUSTERED INDEX PK_MonthlyConsumption
ON Energy.MonthlyConsumption (Year, Month,
DistrictMunicipalityParishCode, VoltageLevel)
ON monthlyConsumptionYearsPartition(Year);

SELECT * FROM sys.partitions as p
WHERE
    p.rows IN (
        SELECT COUNT(*) FROM Energy.MonthlyConsumption
    ) OR p.rows IN (
        SELECT COUNT(*) FROM Energy.MonthlyConsumption
        GROUP BY Year
    )
ORDER BY p.rows DESC
```

To separate a table into multiple partitions it is necessary to create N partitions and to distinguish from them, a partition function, so that we can know in which partition to look for, given a certain record, in this case, we separate them by Year.

Before splitting the table into partitions, it is mandatory to remove the primary key's clustered index, since it organises the records by their primary key inside a single partition. Only after, a new clustered index can be created following the newly created partition function, separating each record into its respective partition.

Task 2

Since the column Year in used in the primary key clustered index, its information is already present in the pointer to the data table. This means that its presence in not necessary in a new index.

It is also important to note that the table is already ordered by year due to the clustered index.

Because the column ActiveEnergy is only used inside a sum statement, it is only important to include their values into the end of the tree instead of sorting the index tree by them.

Initial state

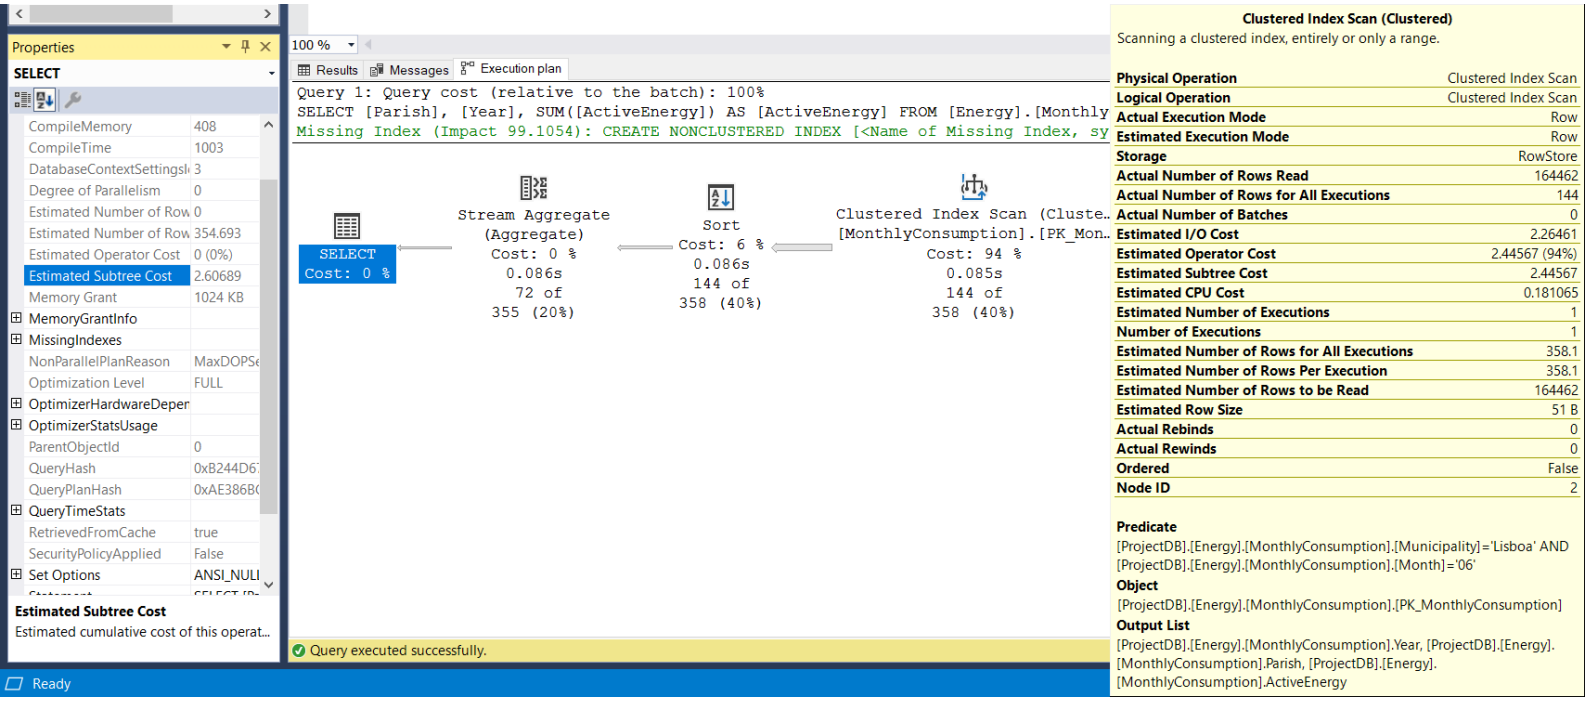

Since the main purpose of an index is to speed up the search of some given records, there is no doubt that an index with Municipality and Month would help greatly, speeding up the look up for the records with the given values, so the first index to try would be:

```
DROP INDEX IF EXISTS IX_Month_Municipality ON Energy.MonthlyConsumption
CREATE NONCLUSTERED INDEX IX_Month_Municipality ON
Energy.MonthlyConsumption (Month, Municipality) INCLUDE (ActiveEnergy)
```

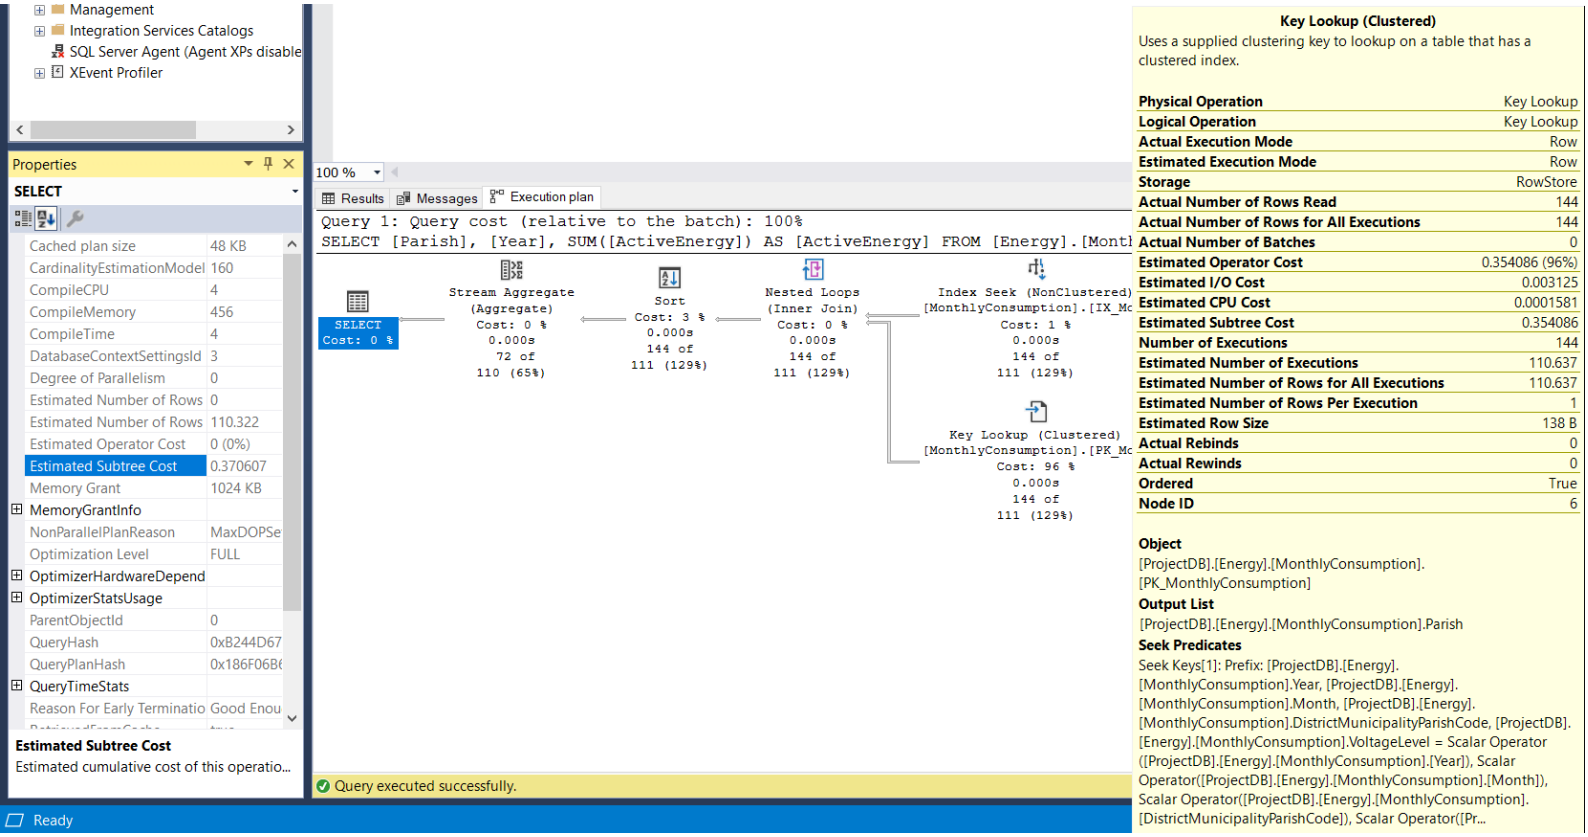

There is still a key lookup being executed using 96% of the total cost. It's only used to get the data for the Parish column. Since it is on a group by and a order by statements, it should go inside the main index clause, but let's include it first, appending its values to the end of the tree.

```
DROP INDEX IF EXISTS IX_Month_Municipality_Recomended ON
Energy.MonthlyConsumption
CREATE NONCLUSTERED INDEX IX_Month_Municipality_Recomended ON
Energy.MonthlyConsumption (Month, Municipality) INCLUDE (Parish,
ActiveEnergy)
```

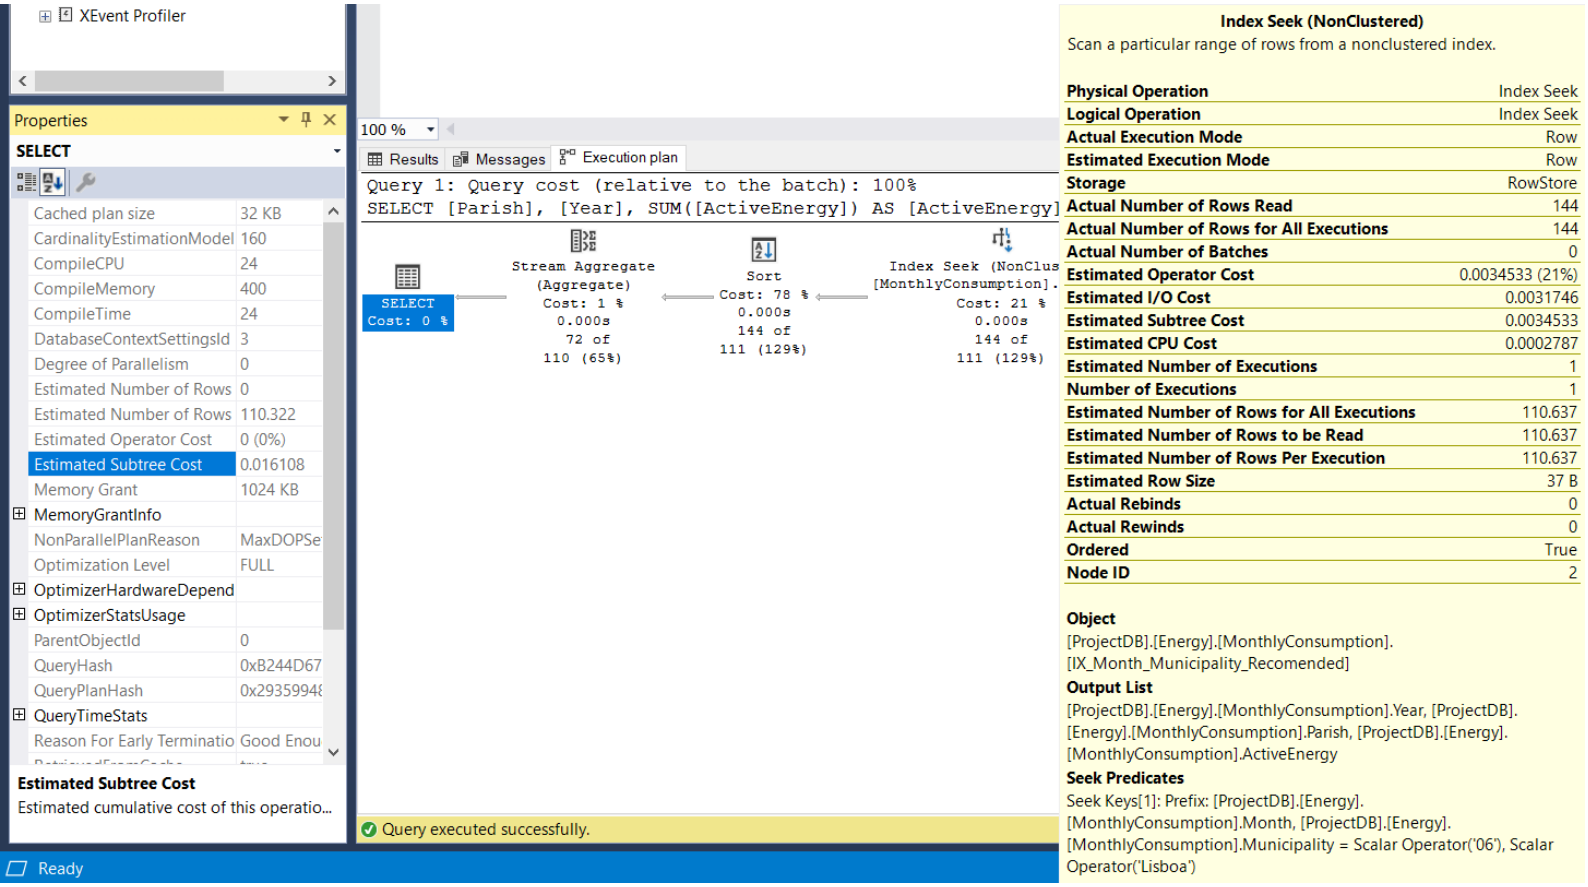

There was a clear benefit on removing that Key Lookup, turning its 0.354 estimated cost into the 0.000345 estimated cost of the hole Index Seek. But now the Sort is the step that is slowing down the most. Moving the Parish column to be used in the tree will automatically sort its values, removing completely the need for a Sort block.

```
DROP INDEX IF EXISTS IX_Month_Municipality_Parish ON
Energy.MonthlyConsumption
CREATE NONCLUSTERED INDEX IX_Month_Municipality_Parish ON
Energy.MonthlyConsumption (Month, Municipality, Parish) INCLUDE
(ActiveEnergy)
```

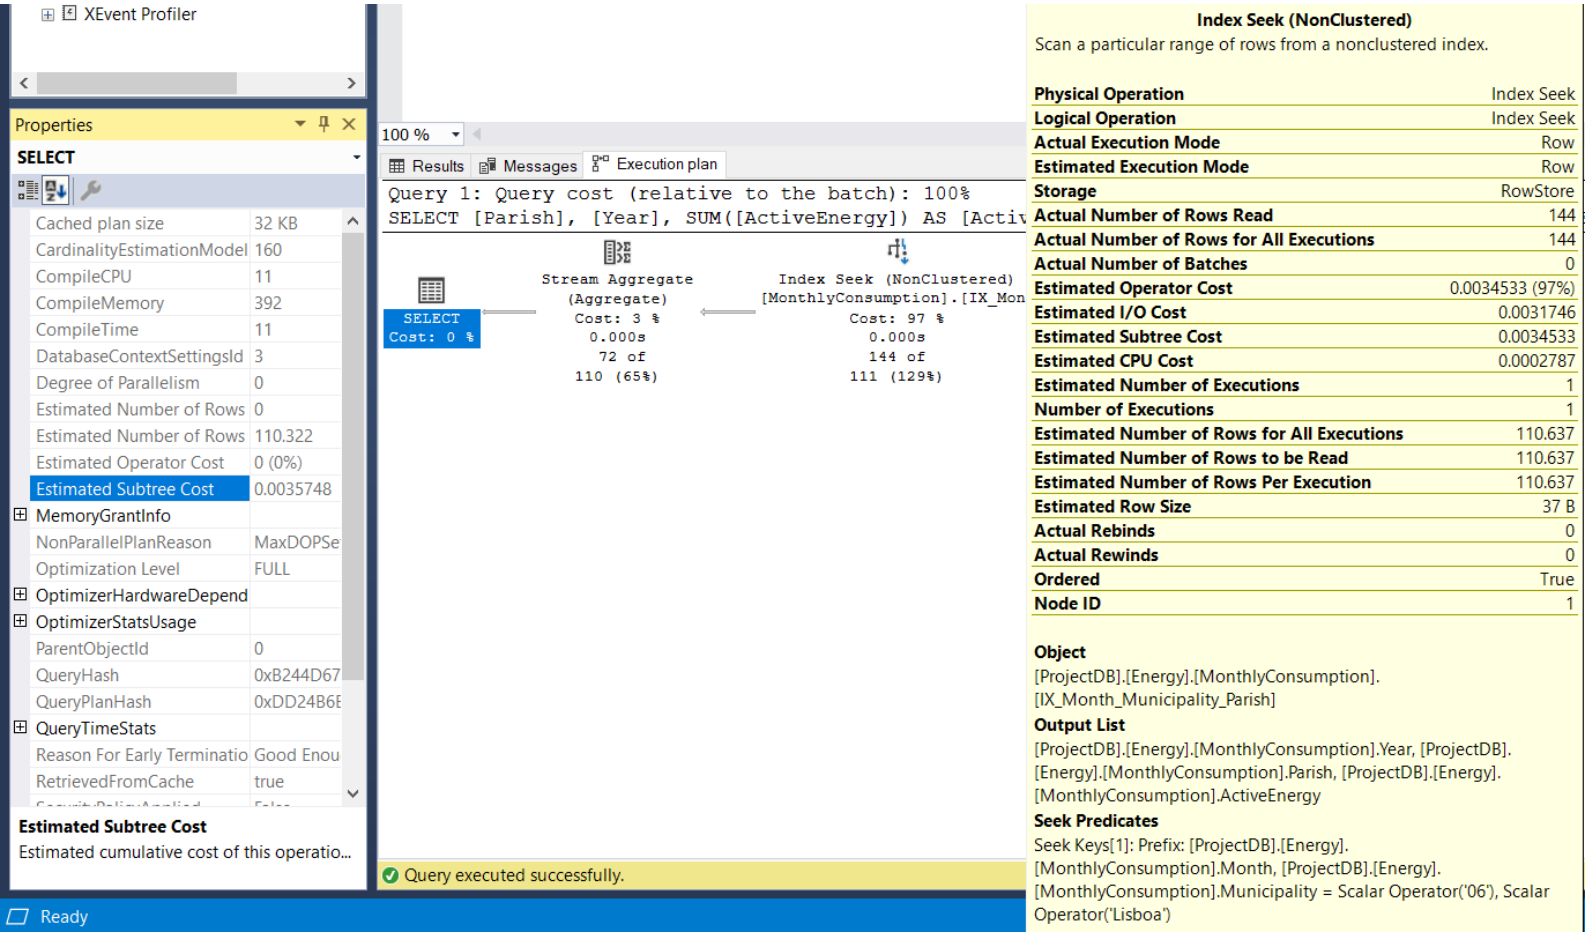

Concluding, the best Index found uses the columns Month, Municipality and Parish, in this order, so after filtering the given Month and Municipality, the Parish is already sorted in the tree, and accesses the Year and ActiveEnergy data from the primary key pointer and the include clause, respectively.

Task 3

The following SQL instructions were added at the end of the given query.

```
OPTION (LOOP JOIN)
```

110389

CardinalityEstimationModel	160
CompileCPU	12
CompileMemory	904
CompileTime	12
DatabaseContextSettingsId	3
Degree of Parallelism	0
Estimated Number of Rows	0
Estimated Number of Rows	8179640
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	110389
Memory Grant	85 MB
MemoryGrantInfo	
NonParallelPlanReason	MaxDOPSe
Optimization Level	FULL
OptimizerHardwareDepend	
OptimizerStatsUsage	
ParentObjectId	0

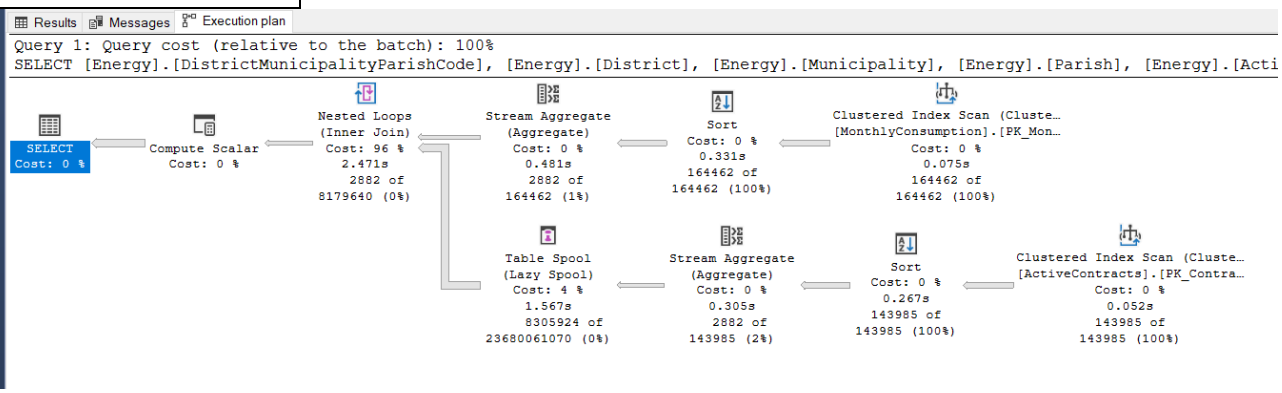


```
OPTION (MERGE JOIN)
```

596.062

CardinalityEstimationModel	160
CompileCPU	17
CompileMemory	864
CompileTime	17
DatabaseContextSettingsId	3
Degree of Parallelism	0
Estimated Number of Rows	0
Estimated Number of Rows	8179640
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	596.062
Memory Grant	636 MB
MemoryGrantInfo	
NonParallelPlanReason	MaxDOPSe
Optimization Level	FULL
OptimizerHardwareDepend	
OptimizerStatsUsage	
ParentObjectId	0

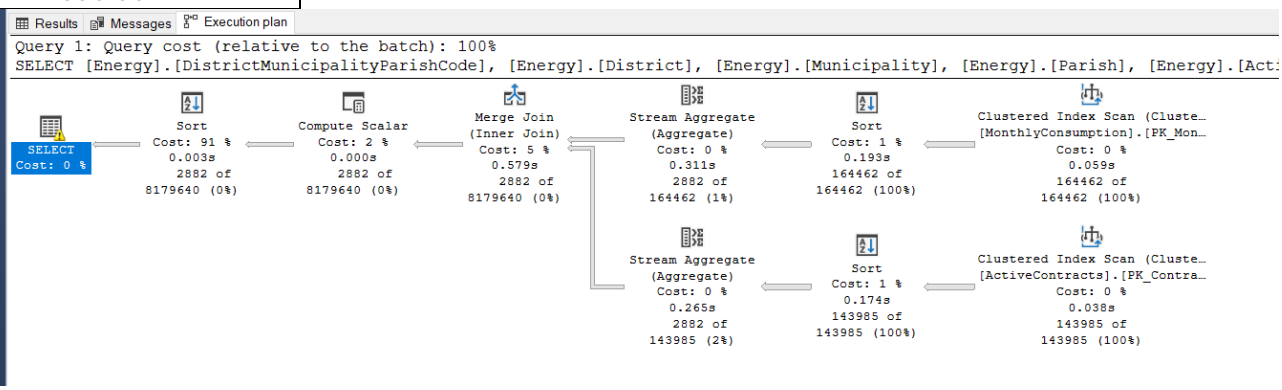


```
OPTION (HASH JOIN)
```

535.201

CardinalityEstimationModel	160
CompileCPU	4
CompileMemory	848
CompileTime	5
DatabaseContextSettingsId	3
Degree of Parallelism	0
Estimated Number of Rows	0
Estimated Number of Rows	8179640
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	535.201
Memory Grant	636 MB
MemoryGrantInfo	
NonParallelPlanReason	MaxDOPSe
Optimization Level	FULL
OptimizerHardwareDepend	
OptimizerStatsUsage	
ParentObjectId	0

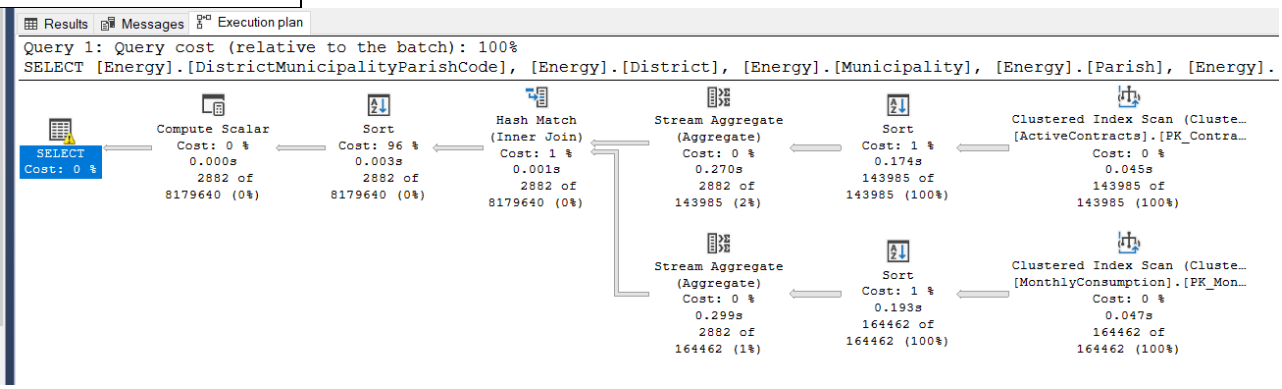

```
OPTION (STREAM AGGREGATE)
```

535.201

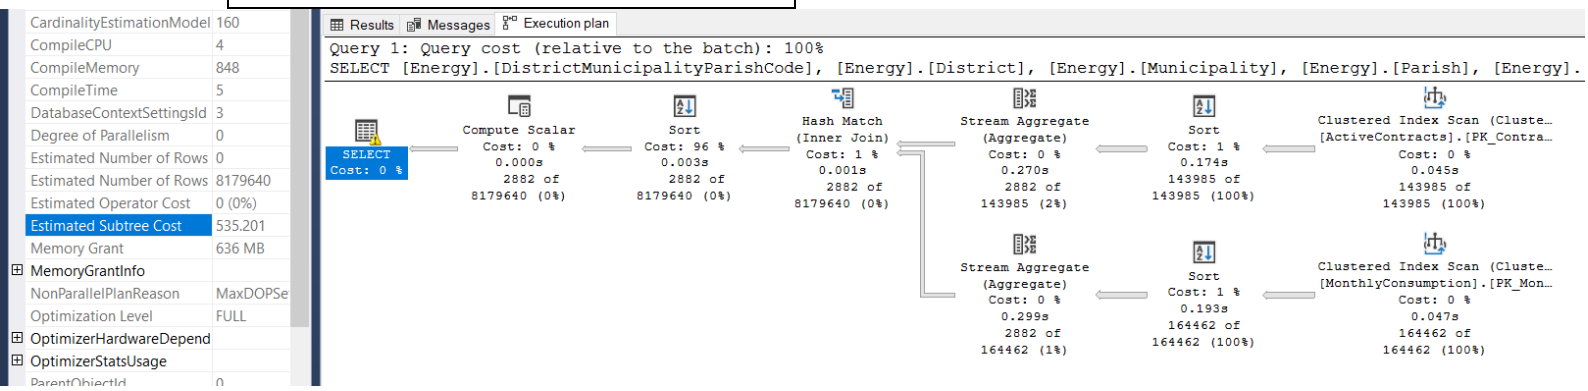


```
OPTION (HASH GROUP)
```

567.669

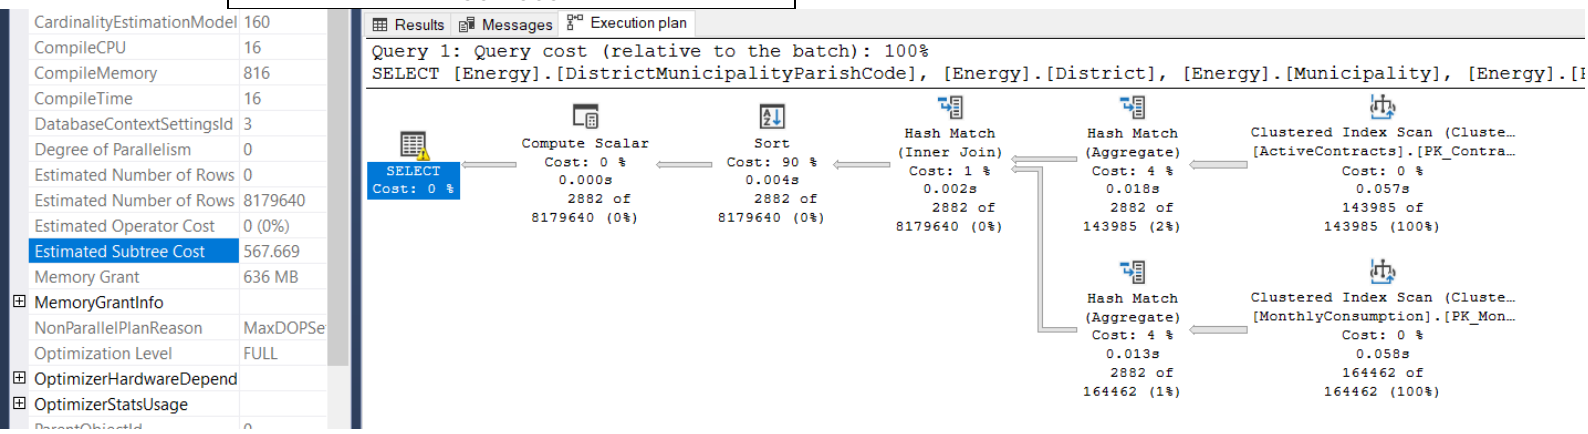

Task 4

```sql
IF OBJECT_ID ('Energy.energy', 'view') IS NOT NULL
   DROP VIEW Energy.energy;
GO
CREATE VIEW Energy.energy WITH SCHEMABINDING AS
    SELECT [DistrictMunicipalityParishCode],
        [District],
        [Municipality],
        [Parish],
        SUM([ActiveEnergy]) AS [ActiveEnergy],
        COUNT_BIG(*) AS COUNT
    FROM [Energy].[MonthlyConsumption]
    GROUP BY [DistrictMunicipalityParishCode],
        [District],
        [Municipality],
        [Parish];
GO
CREATE UNIQUE CLUSTERED INDEX IX_DistrictMunicipalityParishCode
   ON Energy.energy (DistrictMunicipalityParishCode);
GO
```

```sql
IF OBJECT_ID ('Energy.contracts', 'view') IS NOT NULL
   DROP VIEW Energy.contracts;
GO
CREATE VIEW Energy.contracts WITH SCHEMABINDING AS
    SELECT [DistrictMunicipalityParishCode],
        [District],
        [Municipality],
        [Parish],
        SUM([NumberContracts]) AS [NumberContracts],
        COUNT_BIG(*) AS COUNT
    FROM [Energy].[ActiveContracts]
    GROUP BY [DistrictMunicipalityParishCode],
        [District],
        [Municipality],
        [Parish];
GO
CREATE UNIQUE CLUSTERED INDEX IX_DistrictMunicipalityParishCode
   ON Energy.contracts (DistrictMunicipalityParishCode);
GO
```

```sql
SELECT [Energy].[DistrictMunicipalityParishCode],
    [Energy].[District],
    [Energy].[Municipality],
    [Energy].[Parish],
    [Energy].[ActiveEnergy],
    [Contracts].[NumberContracts],
    [Energy].[ActiveEnergy] / [Contracts].[NumberContracts] AS
EnergyPerContract
FROM Energy.energy as [Energy], Energy.contracts as [Contracts]
WHERE [Energy].[DistrictMunicipalityParishCode] =
    [Contracts].[DistrictMunicipalityParishCode]
ORDER BY [Energy].[District],
    [Energy].[Municipality],
    [Energy].[Parish]
```

```
OPTION (LOOP JOIN)
```
27.6459

CardinalityEstimationModel	160
CompileCPU	6
CompileMemory	1160
CompileTime	6
DatabaseContextSettingsId	3
Degree of Parallelism	0
Estimated Number of Rows	0
Estimated Number of Rows	8179640
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	27.6459
Memory Grant	65 MB
MemoryGrantInfo	
NonParallelPlanReason	MaxDOPSe
Optimization Level	FULL
OptimizerHardwareDepend	
OptimizerStatsUsage	
ParentObjectId	0

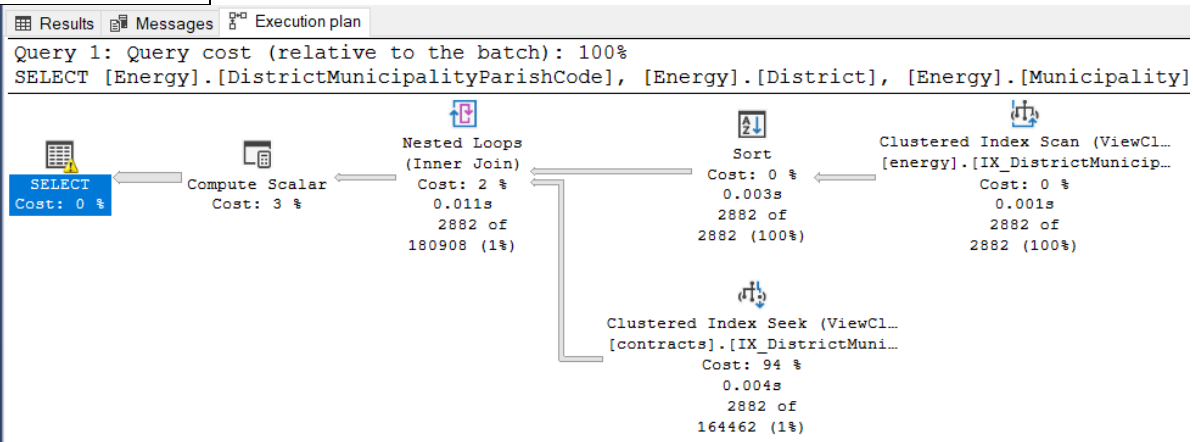


```
OPTION (MERGE JOIN)
```
579.201

CompileCPU	8
CompileMemory	944
CompileTime	8
DatabaseContextSettingsId	3
Degree of Parallelism	0
Estimated Number of Rows	0
Estimated Number of Rows	8179640
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	579.201
Memory Grant	629 MB
MemoryGrantInfo	
NonParallelPlanReason	MaxDOPSe
Optimization Level	FULL
OptimizerHardwareDepend	
OptimizerStatsUsage	
ParentObjectId	0
QueryHash	0x87DB182

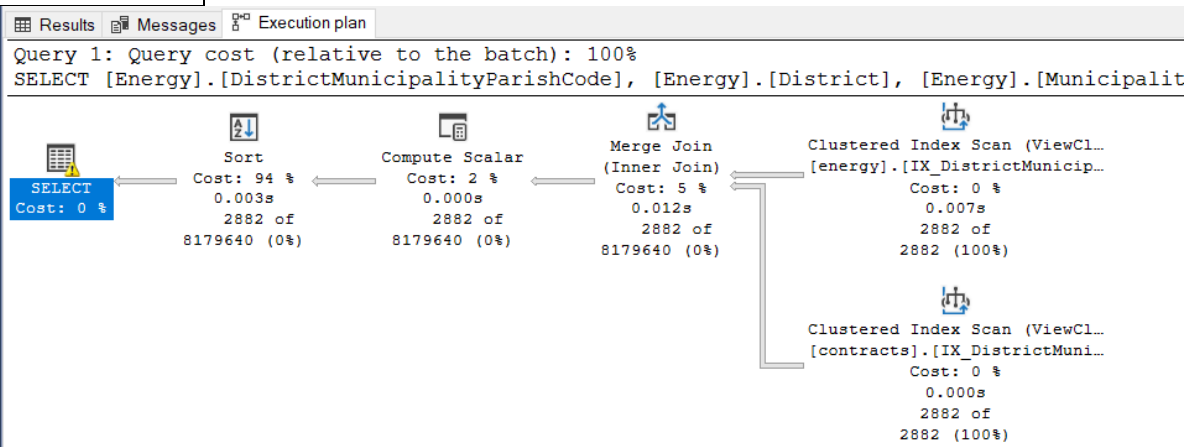


```
OPTION (HASH JOIN)
```
518.119

CardinalityEstimationModel	160
CompileCPU	20
CompileMemory	968
CompileTime	20
DatabaseContextSettingsId	3
Degree of Parallelism	0
Estimated Number of Rows	0
Estimated Number of Rows	8179640
Estimated Operator Cost	0 (0%)
Estimated Subtree Cost	518.119
Memory Grant	629 MB
MemoryGrantInfo	
NonParallelPlanReason	MaxDOPSe
Optimization Level	FULL
OptimizerHardwareDepend	
OptimizerStatsUsage	

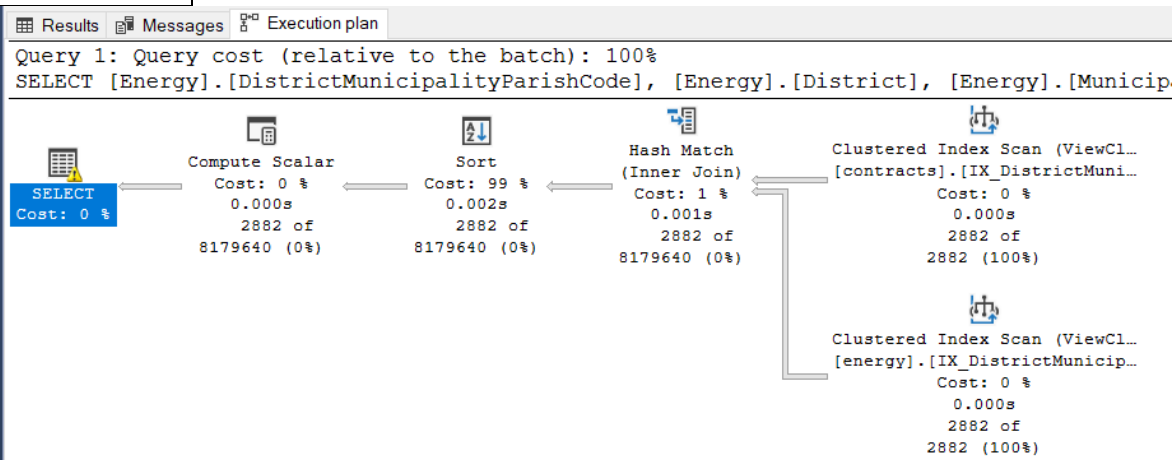

Task 5

In this task in order to get the best results possible run both queries from task 2 and 3 in the same workload in order to get the best results possible that apply to both.

This is our workload:

```
SELECT [Parish], [Year], SUM([ActiveEnergy]) AS [ActiveEnergy]
FROM [Energy].[MonthlyConsumption]
WHERE [Municipality] = 'Lisboa'
AND [Month] = '06'
GROUP BY [Parish], [Year]
ORDER BY [Parish], [Year]

SELECT [Energy].[DistrictMunicipalityParishCode],
    [Energy].[District],
    [Energy].[Municipality],
    [Energy].[Parish],
    [Energy].[ActiveEnergy],
    [Contracts].[NumberContracts],
    [Energy].[ActiveEnergy] / [Contracts].[NumberContracts] AS
EnergyPerContract
FROM (
    SELECT [DistrictMunicipalityParishCode],
        [District],
        [Municipality],
        [Parish],
        SUM([ActiveEnergy]) AS [ActiveEnergy]
    FROM [Energy].[MonthlyConsumption]
    GROUP BY [DistrictMunicipalityParishCode],
        [District],
        [Municipality],
        [Parish]
) AS [Energy], (
    SELECT [DistrictMunicipalityParishCode],
        [District],
        [Municipality],
        [Parish],
    SUM([NumberContracts]) AS [NumberContracts]
    FROM [Energy].[ActiveContracts]
    GROUP BY [DistrictMunicipalityParishCode],
        [District],
        [Municipality],
        [Parish]
) AS [Contracts]
WHERE [Energy].[DistrictMunicipalityParishCode] =
    [Contracts].[DistrictMunicipalityParishCode]
ORDER BY [Energy].[District],
    [Energy].[Municipality],
    [Energy].[Parish]
```

The workload will be run on Database Engine Tuning Advisor with the following settings:

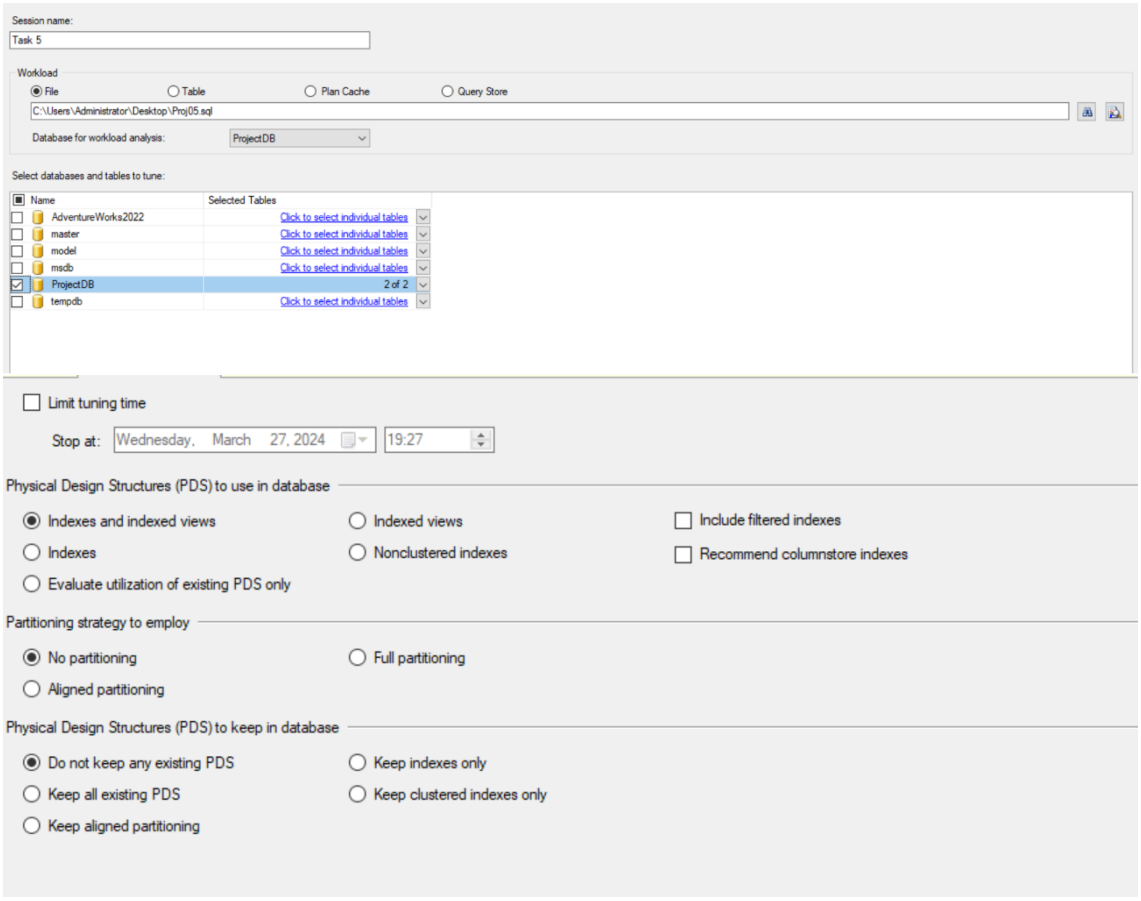

After running the analysis with these settings we get the following recommendations:

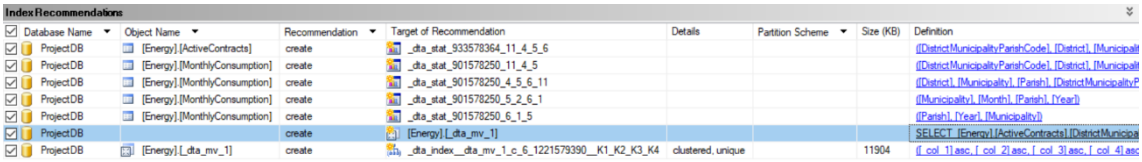

As we can notice the second to last recommendation looks familiar and that is because that is the second materialized view we created in task 4

```
CREATE VIEW [Energy].[_dta_mv_1] WITH SCHEMABINDING
 AS
SELECT  [Energy].[ActiveContracts].[DistrictMunicipalityParishCode] as _col_1,  [Energy].
[ActiveContracts].[District] as _col_2,  [Energy].[ActiveContracts].[Municipality] as _col_3,  [Energy].
[ActiveContracts].[Parish] as _col_4,  SUM([ProjectDB].[Energy].[ActiveContracts].[NumberContracts]) as
_col_5, count_big(*) as _col_6 FROM  [Energy].[ActiveContracts]  GROUP BY  [Energy].
[ActiveContracts].[DistrictMunicipalityParishCode],  [Energy].[ActiveContracts].[District],  [Energy].
[ActiveContracts].[Municipality],  [Energy].[ActiveContracts].[Parish]
```

By applying these recommendations and running a new analysis (with the same settings) we will get new recommendations

Again, as we can notice the third recommendation is the other materialized view we created in task 4

```
CREATE VIEW [Energy].[_dta_mv_1_9987] WITH SCHEMABINDING
 AS
SELECT  [Energy].[MonthlyConsumption].[DistrictMunicipalityParishCode] as _col_1,  [Energy].
[MonthlyConsumption].[District] as _col_2,  [Energy].[MonthlyConsumption].[Municipality] as _col_3,
[Energy].[MonthlyConsumption].[Parish] as _col_4,  count_big(*) as _col_5, SUM([ProjectDB].
[Energy].[MonthlyConsumption].[ActiveEnergy]) as _col_6 FROM  [Energy].[MonthlyConsumption]
GROUP BY  [Energy].[MonthlyConsumption].[DistrictMunicipalityParishCode],  [Energy].
[MonthlyConsumption].[District],  [Energy].[MonthlyConsumption].[Municipality],  [Energy].
[MonthlyConsumption].[Parish]
```

Also, the first recommendation is a new index that we can apply to the query on task 2 to get improved performance.

```
SET ANSI_PADDING ON

CREATE NONCLUSTERED INDEX
[_dta_index_MonthlyConsumption_6_901578250__K2_K5_K6_K1_8]
ON [Energy].[MonthlyConsumption]
(
	[Month] ASC,
	[Municipality] ASC,
	[Parish] ASC,
	[Year] ASC
)
INCLUDE([ActiveEnergy]) WITH (SORT_IN_TEMPDB = OFF,
DROP_EXISTING = OFF, ONLINE = OFF) ON [PRIMARY]
```

After creating this index we can compare the estimated subtree cost before and after

SELECT

Cached plan size	32 KB
Estimated Operator Cost	0 (0%)
Degree of Parallelism	0
Estimated Subtree Cost	2.60689
Memory Grant	1024 KB
Estimated Number of Rows for All Executions	0
Estimated Number of Rows Per Execution	354.693

Statement

```
SELECT [Parish], [Year], SUM([ActiveEnergy]) AS
[ActiveEnergy]
FROM [Energy].[MonthlyConsumption]
WHERE [Municipality] = 'Lisboa'
AND [Month] = '06'
GROUP BY [Parish], [Year]
ORDER BY [Parish], [Year]
```

SELECT

Cached plan size	24 KB
Estimated Operator Cost	0 (0%)
Degree of Parallelism	0
Estimated Subtree Cost	0.0035748
Estimated Number of Rows for All Executions	0
Estimated Number of Rows Per Execution	110.322

Statement

```
SELECT [Parish], [Year], SUM([ActiveEnergy]) AS [ActiveEnergy]
FROM [Energy].[MonthlyConsumption]
WHERE [Municipality] = 'Lisboa'
AND [Month] = '06'
GROUP BY [Parish], [Year]
ORDER BY [Parish], [Year]
```